

HAIYANG SHI

🏠 2015 Neil Avenue, Columbus, OH 43210

✉ shi.876@buckeyemail.osu.edu

☎ 614-607-4016

SUMMARY OF QUALIFICATIONS

- 7+ years of research and work experience in designing and developing high-performance and scalable systems
- In-depth knowledge of high-performance interconnects (e.g., Remote Direct Memory Access (RDMA)) and protocols, Erasure Coding (EC), In-Network Computing, and distributed systems
- Skilled in programming languages: C/C++, Java, Python, SQL, and Shell scripting
- Years of experience in using message queues, databases (e.g., MySQL, HBase, TitanDB), and Big Data processing frameworks (e.g., Hadoop, Spark)
- Strong problem-solving and communication skills
- Self-motivated and efficient in cross-group and cross-culture collaboration

EDUCATION

The Ohio State University

Ph.D. in Computer Science and Engineering. Advisor: Prof. Xiaoyi Lu

Columbus, OH

Aug 2016 – Dec 2020

Tianjin University

B.E. in Computer Science and Technology

Tianjin, China

Aug 2008 – Jul 2012

EXPERIENCE

Graduate Research Assistant

The Ohio State University, Columbus, OH

Aug 2016 - Dec 2020

- **High-Performance Erasure Coding Schemes for Next-Generation Storage Systems:** Research project; Multiple papers get published in top conferences (C++/C/Java)
 - Leading the research and implementation of involved EC libraries and schemes
 - Designed and developed a multi-rail EC library leveraging heterogeneous hardware, a tripartite graph based EC scheme taking advantage of EC NIC (Network Interface Card) offloading, and a set of in-network EC primitives to accelerate EC computation in distributed storage systems
 - Invented a set of micro-benchmarks to evaluate the performance of various EC schemes
- **Fast and Scalable Remote Procedure Call (RPC) Frameworks:** Accelerated Apache Thrift and brpc for modern data centers (C++)
 - Leading the design and development of RDMA-oriented RPC frameworks
 - Characterizing and applying networking optimizations for fast message delivery
 - Designed and implemented efficient flow control and communication protocols
 - Participated in benchmarking and testing the frameworks in several data centers
 - RPC frameworks have been exploited in an enterprise database system and demonstrate up to 20% speedup on 1TB TPC-H benchmark
- **Big Data Processing Frameworks over High-Performance Networks:** High-performance derivative of Apache Hadoop accelerated by RDMA technologies (C, Java)
 - Redesigned and implemented RDMA-accelerated I/O pipelines for both replication and erasure coding storage schemes in HDFS 3.x
 - Involved in benchmarking and testing related software stacks
 - Benchmark results show that the enhanced Hadoop improves overall throughput by up to $3.1\times$
 - Open-sourced as a part of OSU HiBD project (<http://hibd.cse.ohio-state.edu>)

Research Intern

Alibaba, Seattle, WA. Mentor: Liwei Peng

May 2019 – Aug 2019

- Benchmarked distributed training of Transformer and BERT models with varied configurations to study the impact of RDMA interconnects and all-reduce algorithms
- Developed tools to analyze the bottleneck of distributed machine learning

Software Engineer

MiningLamp, Beijing, China

May 2015 – Jul 2016

- **Scopa**: A property graph based analysis system presenting data as meaningfully defined objects and relationships: people, places, events, and the connections between them. (*Java, HBase, Spark, TitanDB*)
 - One of the major engineers in designing and developing the system
 - Developed a graph traversal engine for the underlying graph database - TitanDB (or JanusGraph)

Sina Weibo (China's Twitter), Beijing, China

Jul 2012 – Apr 2014

- Participated in designing and developing several subsystems for managing profiles and posts of hundreds of millions of users (*PHP, Message Queue, Memcached, MySQL*)

SELECTED PUBLICATIONS

[SC'20] INEC: Fast and Coherent In-Network Erasure Coding, *H. Shi, and X. Lu*, In Proceedings of the 33rd International Conference for High Performance Computing, Networking, Storage and Analysis (SC), 2020.

[SC'19] TriEC: Tripartite Graph Based Erasure Coding NIC Offload, *H. Shi, and X. Lu*, In Proceedings of the 32nd International Conference for High Performance Computing, Networking, Storage and Analysis (SC), 2019. **Best Student Paper Finalist**.

[HPDC'19] UMR-EC: A Unified and Multi-Rail Erasure Coding Library for High-Performance Distributed Storage Systems, *H. Shi, X. Lu, D. Shankar, and D. K. Panda*, In Proceedings of the 28th ACM International Symposium on High-Performance Parallel and Distributed Computing (HPDC), 2019.

[Bench'18] EC-Bench: Benchmarking Onload and Offload Erasure Coders on Modern Hardware Architectures, *H. Shi, X. Lu, and D. K. Panda*, In Proceedings of International Symposium on Benchmarking, Measuring, and Optimizing (Bench), 2018. **Best Paper**.

[HotI'17] Characterizing Deep Learning over Big Data (DLoBD) Stacks on RDMA-capable Networks, *X. Lu, H. Shi, H. Javed, R. Biswas, and D. K. Panda*, In Proceedings of the 25th Annual Symposium on High-Performance Interconnects (HotI), 2017.

AWARDS

- SC'19 Best Student Paper Finalist
- Bench'18 Best Paper Award

PROFESSIONAL INVOLVEMENT

- Publicity chair of HPBDC'20
- Reviewer of TPDS'19
- External/Sub reviewer of TPDS Special Section on AI/ML/DL 2020, HiPC'20, Bench'20, UCC'20, CCGrid'21